

Lecture 5.0

Shortest Paths

Department of Computer Science
Hofstra University

Lecture Goals

- In this lecture we study **shortest-paths** problems. We begin by analyzing some basic properties of shortest paths and a generic algorithm for the problem.
- We introduce and analyze **Dijkstra's algorithm** for shortest-paths problems with nonnegative weights.
- We conclude with the **Bellman–Ford** algorithm for edge-weighted digraphs with no negative cycles.

Shortest Paths in an Edge-weighted Digraph

Given an edge-weighted digraph, find the shortest path from s to t .

edge-weighted digraph

4→5	0.35
5→4	0.35
4→7	0.37
5→7	0.28
7→5	0.28
5→1	0.32
0→4	0.38
0→2	0.26
7→3	0.39
1→3	0.29
2→7	0.34
6→2	0.40
3→6	0.52
6→0	0.58
6→4	0.93



shortest path from 0 to 6

0→2	0.26
2→7	0.34
7→3	0.39
3→6	0.52

Variants

❖ Which vertices?

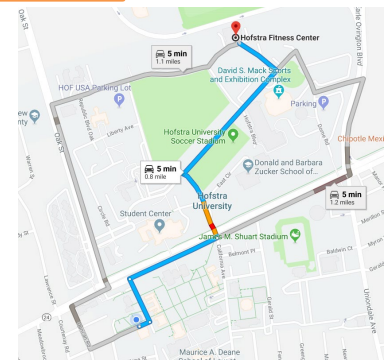
- Single source: from one vertex s to every other vertex.
- Source-sink: from one vertex s to another t .
- All pairs: between all pairs of vertices.

❖ Nonnegative weights?

❖ Cycles?

- Negative cycles.

Can we use BFS?



Simplifying assumption: Each vertex is reachable from s .

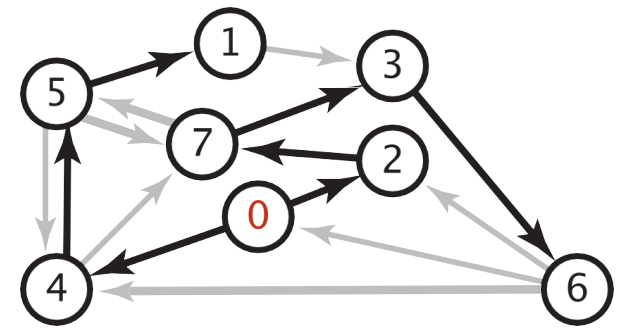
Data Structures for Single-source Shortest Paths

Goal. Find the shortest path from s to every other vertex.

Observation. A shortest-paths tree (SPT) solution exists.

Consequence. Can represent the SPT with two vertex-indexed arrays:

- $\text{distTo}[v]$ is length of shortest path from s to v .
- $\text{edgeTo}[v]$ is last edge on shortest path from s to v .



shortest-paths tree from 0

	edgeTo[]	distTo[]
0	null	0
1	5->1 0.32	1.05
2	0->2 0.26	0.26
3	7->3 0.37	0.97
4	0->4 0.38	0.38
5	4->5 0.35	0.73
6	3->6 0.52	1.49
7	2->7 0.34	0.60

parent-link representation


```
public double distTo(int v)
{ return distTo[v]; }

public Iterable<DirectedEdge> pathTo(int v)
{
    Stack<DirectedEdge> path = new Stack<DirectedEdge>();
    for (DirectedEdge e = edgeTo[v]; e != null; e = edgeTo[e.from()])
        path.push(e);
    return path;
}
```

Edge Relaxation

Relax edge $e = v \rightarrow w$. (basic of building SPT)

- distTo[v] is length of shortest **known** path from s to v.
- distTo[w] is length of shortest **known** path from s to w.
- edgeTo[w] is last edge on shortest **known** path from s to w.
- If $e = v \rightarrow w$ gives shorter path to w through v, update distTo[w] and edgeTo[w].
 - distTo[w] = min(distTo[w], distTo[v] + e.weight());



```
private void relax(DirectedEdge e)
{
    int v = e.from(), w = e.to();
    if (distTo[w] > distTo[v] + e.weight())
    {
        distTo[w] = distTo[v] + e.weight();
        edgeTo[w] = e;
    }
}
```

OLD distTo[w] = 7.2 > distTo[v] + e.weight() = 3.1 + 1.3 = 4.4
 NEW distTo[w] ← distTo[v] + e.weight() = 4.4

Generic Shortest-paths Algorithm

Generic algorithm (to compute SPT from s)

For each vertex v : $\text{distTo}[v] = \infty$.

For each vertex v : $\text{edgeTo}[v] = \text{null}$.

$\text{distTo}[s] = 0$.

Repeat until done:

- Relax any edge.

Proposition. Generic algorithm computes SPT (if it exists) from s .

Pf.

- Throughout algorithm, $\text{distTo}[v]$ is the length of a simple path from s to v (and $\text{edgeTo}[v]$ is last edge on path).
- Each successful relaxation decreases $\text{distTo}[v]$ for some v .
- The entry $\text{distTo}[v]$ can decrease at most a finite number of times.

Efficient implementations. How to choose which edge to relax?

- Ex 1. Dijkstra's algorithm. (**nonnegative weights, directed cycles**).
- Ex 2. Bellman–Ford algorithm. (**no negative cycles**).



Dijkstra's Algorithm

- Initialization:
 - Set the distance to the source node as 0 and to all other nodes as infinity.
 - Mark all nodes as unvisited and store them in a priority queue.
- Main Loop:
 - Extract the **unvisited node** with **the smallest known distance** from the queue.
 - For each neighboring node, calculate its tentative distance through the current node. **If this distance is smaller than the previously recorded distance, update it with relaxation.**
 - Mark the current node as visited once all its neighbors are processed.
- Termination:
 - The algorithm continues until all reachable nodes are visited, or until the shortest path to a specific destination is found.

Dijkstra's Algorithm: Correctness Proof

Proposition. Dijkstra's algorithm computes a SPT in any edge-weighted digraph with nonnegative weights.

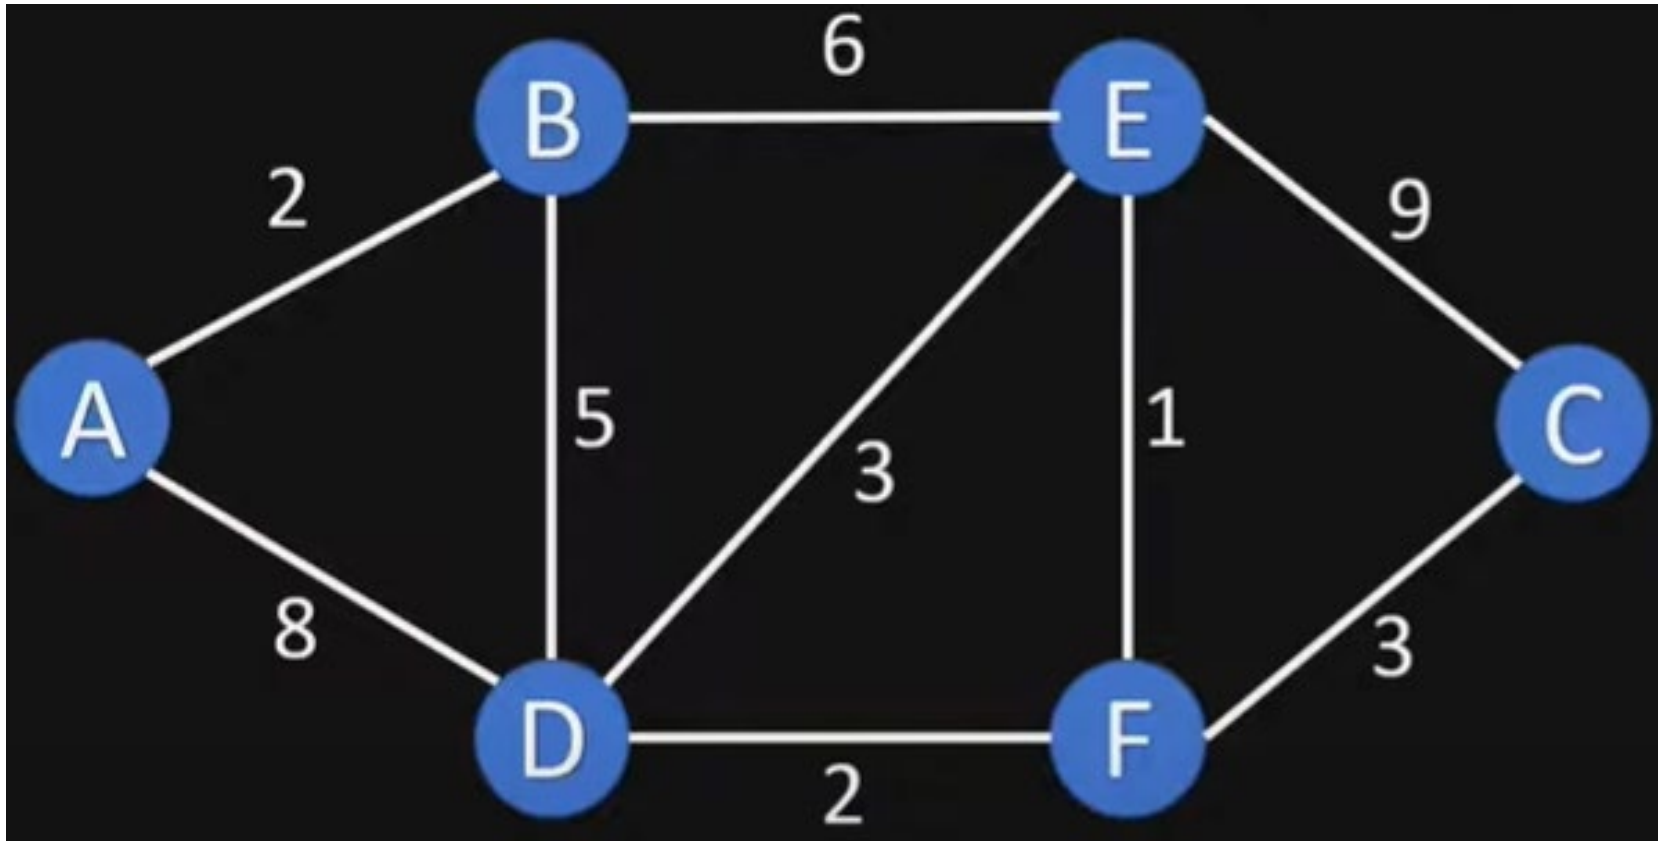
Pf.

- Each edge $e = v \rightarrow w$ is relaxed exactly once (when v is relaxed),
 - leaving $\text{distTo}[w] \leq \text{distTo}[v] + e.\text{weight}()$.
- Inequality holds until algorithm terminates because:
 - $\text{distTo}[w]$ cannot increase  $\text{distTo}[\]$ values are monotone decreasing
 - $\text{distTo}[v]$ will not change  we choose lowest $\text{distTo}[\]$ value at each step (and edge weights are nonnegative)
- Thus, upon termination, shortest-paths optimality conditions hold.

Video Tutorial


- The following lecture slides are based on this video:
- Dijkstras Shortest Path Algorithm Explained | With Example | Graph Theory
 - <https://www.youtube.com/watch?v=bZkzH5x0SKU>

Example Graph



Initialize

2. Assign to all nodes a tentative distance value



Visited Nodes: []

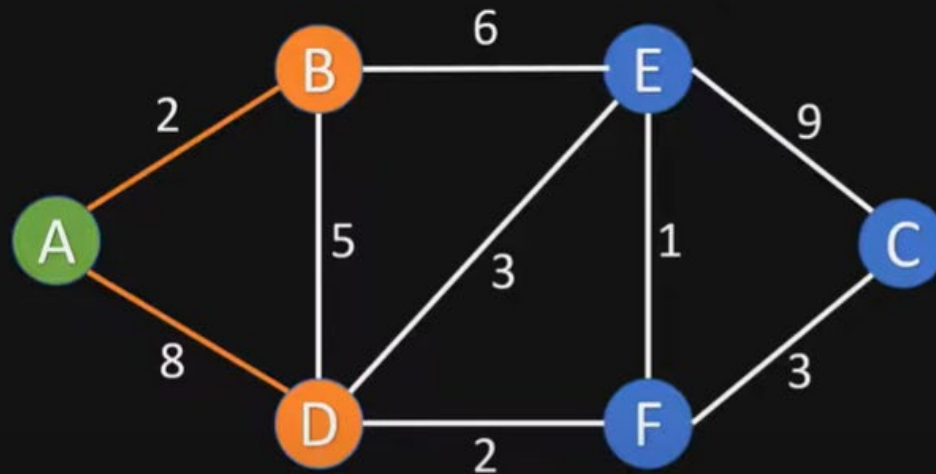
Unvisited Nodes: [A, B, C, D, E, F]

Node	Shortest Distance	Previous Node
A	0	
B	∞	
C	∞	
D	∞	
E	∞	
F	∞	

Visit Node A

3. For the current node calculate the distance to all unvisited neighbours

3.1. Update shortest distance, if new distance is shorter than old distance



Node	Shortest Distance	Previous Node
A	0	
B	2	A
C	∞	
D	8	A
E	∞	
F	∞	

Visited Nodes: []

Unvisited Nodes: [A, B, C, D, E, F]

OLD $\text{distTo}[B] = \infty > \text{distTo}[A] + e[A][B].\text{weight}() = 0 + 2 = 2$

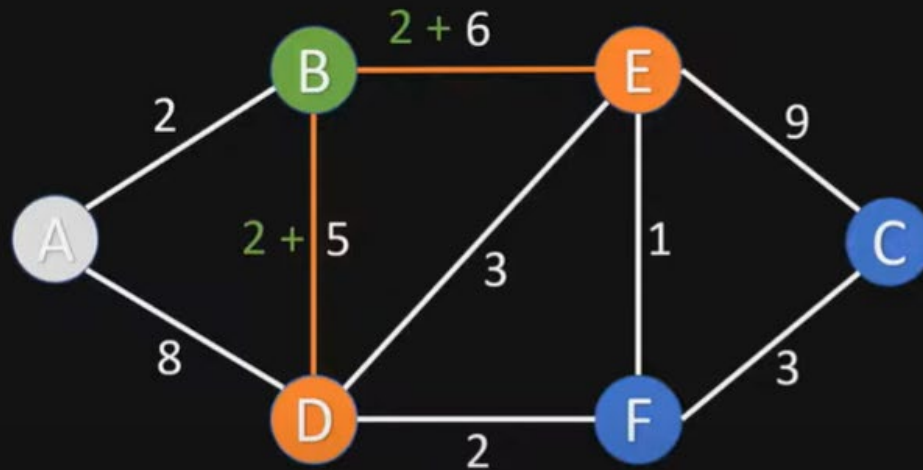
NEW $\text{distTo}[B] \leftarrow \text{distTo}[A] + e[A][B].\text{weight}() = 2$

OLD $\text{distTo}[D] = \infty > \text{distTo}[A] + e[A][D].\text{weight}() = 0 + 8 = 8$

NEW $\text{distTo}[D] \leftarrow \text{distTo}[A] + e[A][D].\text{weight}() = 8$

Visit Node B

3. For the current node calculate the distance to all unvisited neighbours
 3.1. Update shortest distance, if new distance is shorter than old distance



Visited Nodes: [A]

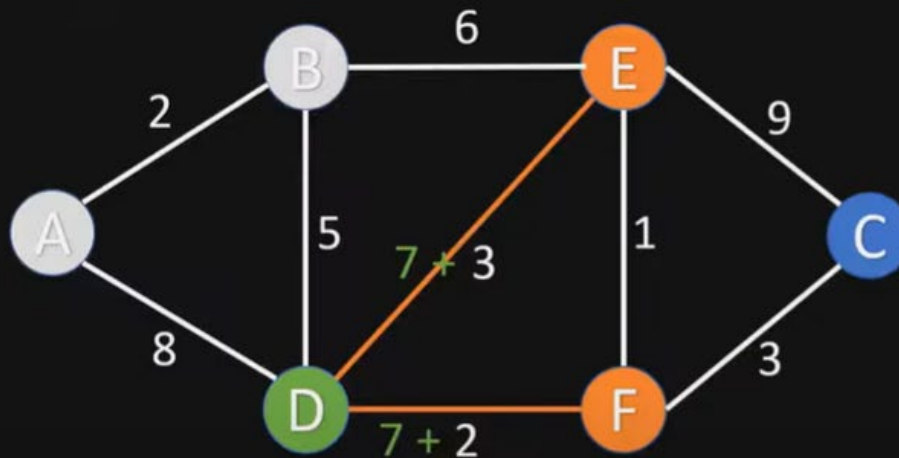
Unvisited Nodes: [B, C, D, E, F]

Node	Shortest Distance	Previous Node
A	0	
B	2	A
C	∞	
D	7	B
E	8	B
F	∞	

OLD $\text{distTo}[D] = 8 > \text{distTo}[B] + e[B][D].\text{weight}() = 2 + 5 = 7$
 NEW $\text{distTo}[D] \leftarrow \text{distTo}[B] + e[B][D].\text{weight}() = 7$
 OLD $\text{distTo}[E] = \infty > \text{distTo}[B] + e[B][E].\text{weight}() = 2 + 6 = 8$
 NEW $\text{distTo}[E] \leftarrow \text{distTo}[B] + e[B][E].\text{weight}() = 8$

Visit Node D

3. For the current node calculate the distance to all unvisited neighbours
 3.1. Update shortest distance, if new distance is shorter than old distance




Visited Nodes: [A, B] Unvisited Nodes: [C, D, E, F]

Node	Shortest Distance	Previous Node
A	0	
B	2	A
C	∞	
D	7	B
E	8	B
F	9	D

$\text{OLD distTo}[E] = 8 < \text{distTo}[D] + e[D][E].\text{weight}() = 7+3 = 10$
 No update, $\text{distTo}[E]$ stays 8
 $\text{OLD distTo}[F] = \infty > \text{distTo}[D] + e[D][F].\text{weight}() = 7+2 = 9$
 $\text{NEW distTo}[F] \leftarrow \text{distTo}[D] + e[D][F].\text{weight}() = 9$

Visit Node E

3. For the current node calculate the distance to all unvisited neighbours
- 3.1. Update shortest distance, if new distance is shorter than old distance



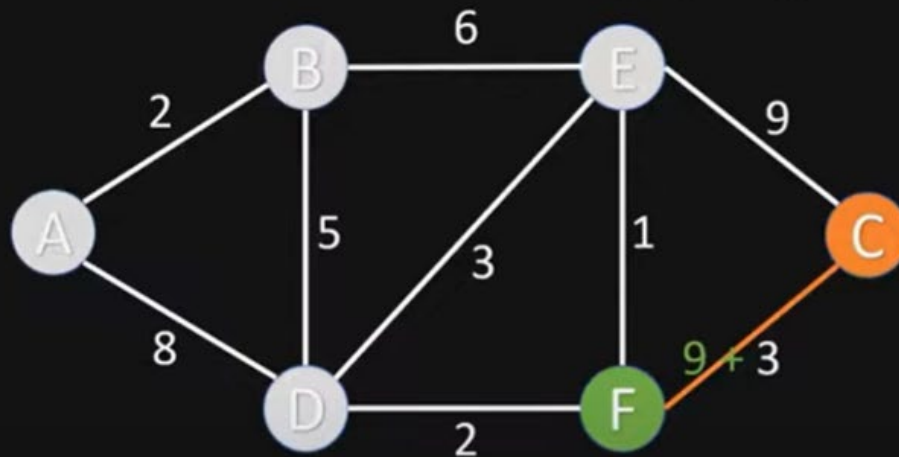
Visited Nodes: [A, B, D] Unvisited Nodes: [C, E, F]

Node	Shortest Distance	Previous Node
A	0	
B	2	A
C	17	E
D	7	B
E	8	B
F	9	D

$\text{OLD distTo}[C] = \infty > \text{distTo}[E] + e[E][C].\text{weight}() = 8 + 9 = 17$
 $\text{NEW distTo}[C] \leftarrow \text{distTo}[E] + e[E][C].\text{weight}() = 17$
 $\text{OLD distTo}[F] = 9 = \text{distTo}[E] + e[E][F].\text{weight}() = 8 + 1 = 9$
 No update, $\text{distTo}[F]$ stays 9

Visit Node F

3. For the current node calculate the distance to all unvisited neighbours
 3.1. Update shortest distance, if new distance is shorter than old distance



Visited Nodes: [A, B, D, E] Unvisited Nodes: [C, F]


Node	Shortest Distance	Previous Node
A	0	
B	2	A
C	12	F
D	7	B
E	8	B
F	9	D

OLD $\text{distTo}[C] = 17 > \text{distTo}[F] + e[F][C].\text{weight}() = 9 + 3 = 12$
 NEW $\text{distTo}[C] \leftarrow \text{distTo}[F] + e[F][C].\text{weight}() = 12$

End of Algorithm

- Table contains the shortest distance to each node N from the source node A, and its previous node in the shortest path

4. Mark current node as visited



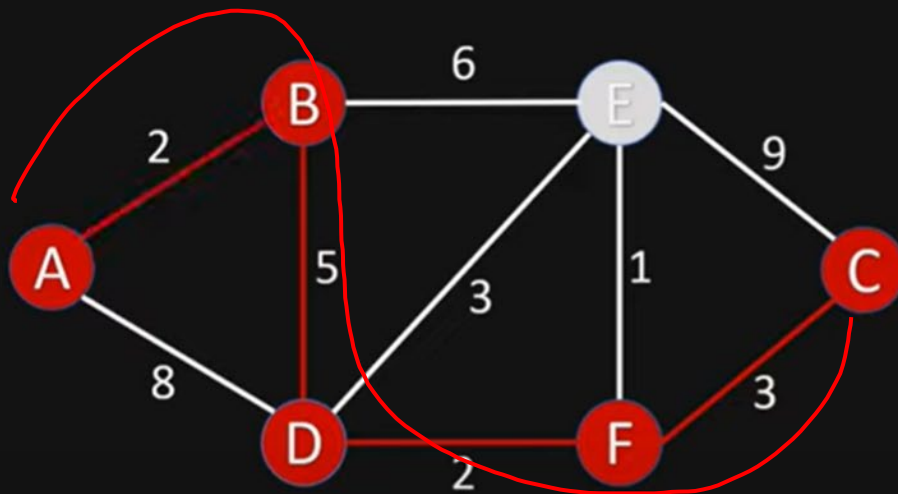
Visited Nodes: [A, B, D, E, F, C] Unvisited Nodes: []

Node	Shortest Distance	Previous Node
A	0	
B	2	A
C	12	F
D	7	B
E	8	B
F	9	D

Getting the Shortest Path from A to C

- C's previous node is F; F's previous node is D; D's previous node is B; B's previous node is A
- Shortest Path from A to C is ABDFC

Get shortest path from A to C

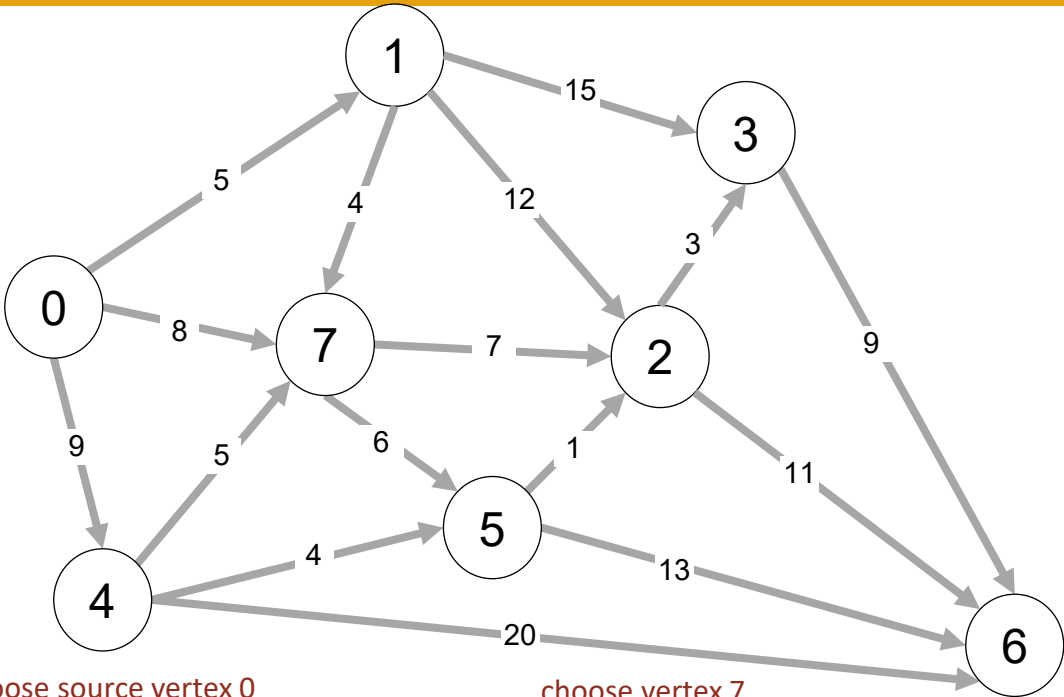


Node	Shortest Distance	Previous Node
A	0	
B	2	A
C	12	F
D	7	B
E	8	B
F	9	D

Dijkstra's Algorithm Example 2

- Consider vertices in increasing order of distance from s
 - (non-tree vertex with the lowest distTo[] value).
- Add vertex to tree and relax all edges pointing from that vertex.

choose vertex 5
 relax all edges adjacent from 5
 choose vertex 2
 relax all edges adjacent from 2
 choose vertex 3
 relax all edges adjacent from 3
 choose vertex 6
 relax all edges adjacent from 6




choose source vertex 0
 relax all edges adjacent from 0
 choose vertex 1
 relax all edges adjacent from 1

choose vertex 7
 relax all edges adjacent from 7
 choose vertex 4
 relax all edges adjacent from 4

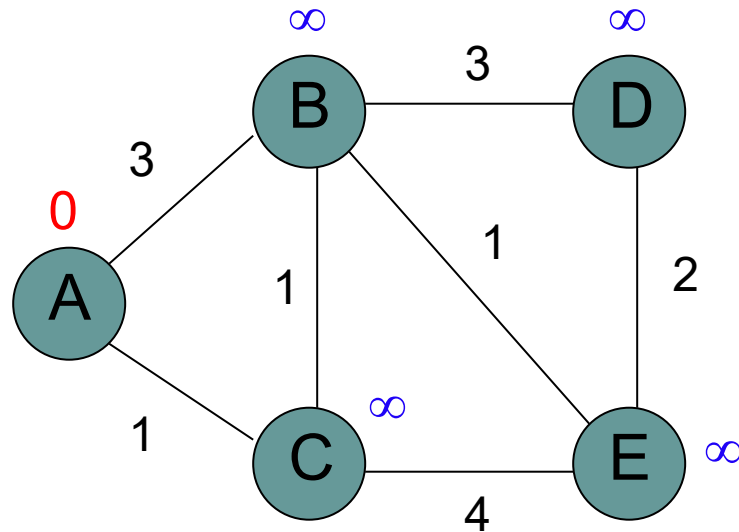
v distTo[]			
→ 0	∞	0	
→ 1	∞	5	
→ 2	∞	17	15 14
→ 3	∞	20	17
→ 4	∞	9	
→ 5	∞	14	13
→ 6	∞	29	26 25
→ 7	∞	8	

v edgeTo[]			
0	-		
1	-	0	
2	-	1	7 5
3	-	1	2
4	-	0	
5	-	7	4
6	-	4	5 2
7	-	0	

Dijkstra's Algorithm Example 3

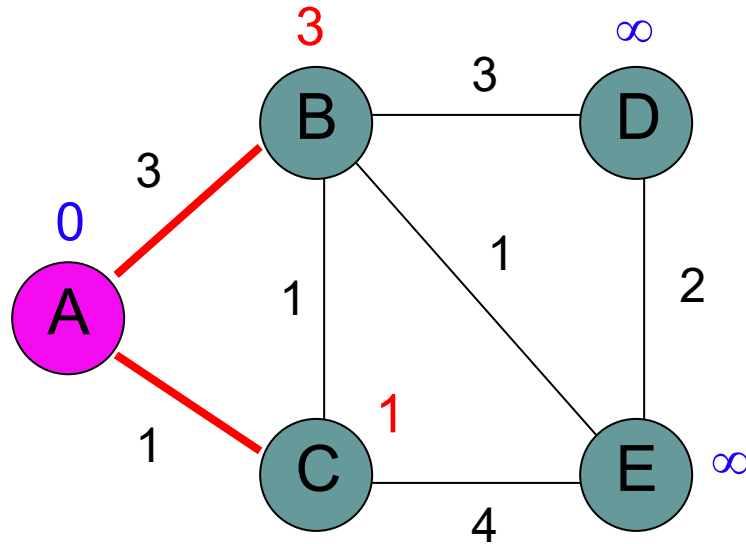


Initialize




N	SD	PN
A	0	
B	∞	
C	∞	
D	∞	
E	∞	

Visit Node A



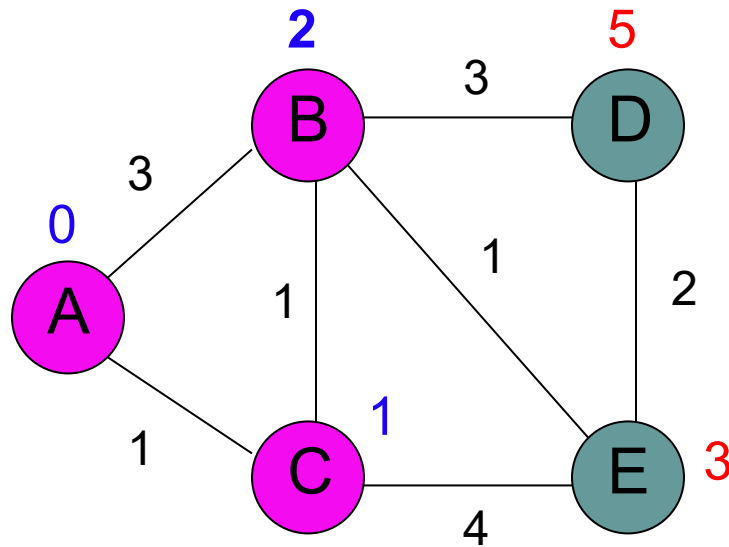
N	SD	PN
A	0	
B	3	A
C	1	A
D	∞	
E	∞	

Visit Node C




N	SD	PN
A	0	
B	2	C
C	1	A
D	∞	
E	5	C

Visit Node B



N	SD	PN
A	0	
B	2	C
C	1	A
D	5	B
E	3	B

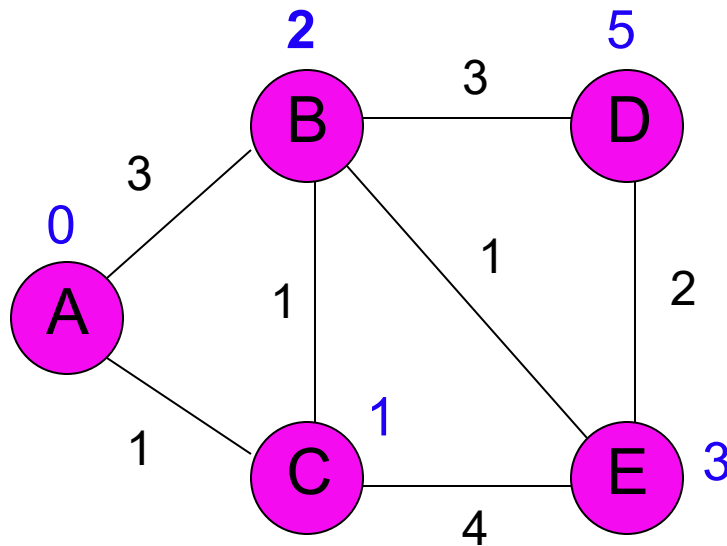
Visit Node E



N	SD	PN
A	0	
B	2	C
C	1	A
D	5	B
E	3	B

Nothing changes

Visit Node D

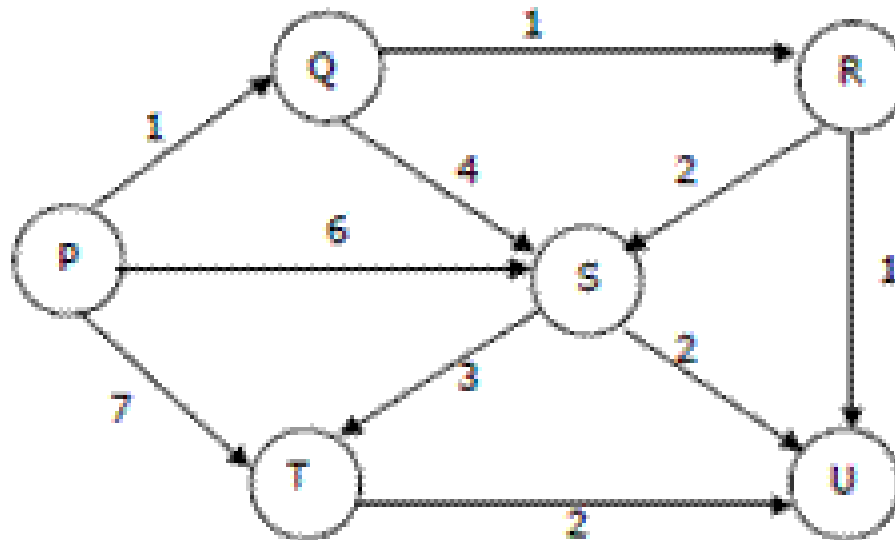


N	SD	PN
A	0	
B	2	C
C	1	A
D	5	B
E	3	B

Nothing changes


Quiz: Dijkstra's Algorithm

- Suppose we run Dijkstra's single source shortest-path algorithm on the following edge weighted directed graph with vertex P as the source. In what order do the nodes get included into the set of vertices for which the shortest path distances are finalized?
- ANS: P, Q, R, U, S, T



SD: Shortest Distance

PN: Previous Node



N	SD	PN
P	0	
Q	∞	
R	∞	
S	∞	
T	∞	
U	∞	

Visit P
→

N	SD	PN
P	0	
Q	1	P
R	∞	
S	6	P
T	7	P
U	∞	

Visit Q
→

N	SD	PN
P	0	
Q	1	P
R	2	Q
S	5	Q
T	7	P
U	∞	

Visit R
→

N	SD	PN
P	0	
Q	1	P
R	2	Q
S	4	Q
T	7	P
U	3	R

← Visit U (nothing changes)

N	SD	PN
P	0	
Q	1	P
R	2	Q
S	4	Q
T	7	P
U	3	R

Visit S
(nothing changes)
→

N	SD	PN
P	0	
Q	1	P
R	2	Q
S	4	Q
T	7	P
U	3	R

Visit T
(nothing changes)
→

N	SD	PN
P	0	
Q	1	P
R	2	Q
S	4	Q
T	7	P
U	3	R

Finished
→

N	SD	PN
P	0	
Q	1	P
R	2	Q
S	4	Q
T	7	P
U	3	R

Bellman-Ford Algorithm

- Initialize distance array `distTo[]` for each vertex `v` as **`distTo[v] = ∞`**, and **`distTo[s] = 0`** to source vertex `s`.
- Relax all **edges** $|V|-1$ times.

```
private void relax(DirectedEdge e)
{
    int v = e.from(), w = e.to();
    if (distTo[w] > distTo[v] + e.weight())
    {
        distTo[w] = distTo[v] +
            e.weight();
        edgeTo[w] = e;
    }
}
```

Recall:

Generic algorithm (to compute SPT from `s`)

For each vertex `v`: `distTo[v] = ∞`.

For each vertex `v`: `edgeTo[v] = null`.

`distTo[s] = 0`.

Repeat until done:

- Relax any edge.

Bellman-Ford algorithm

For each vertex `v`: `distTo[v] = ∞`.

For each vertex `v`: `edgeTo[v] = null`.

`distTo[s] = 0`.

Repeat $|V| - 1$ times:

- Relax each edge.

Bellman-Ford Algorithm Proof of Correctness

- Relaxing edges $|V|-1$ times in the Bellman-Ford Algorithm guarantees that the algorithm has explored all possible paths of length up to $|V|-1$, which is the maximum possible length of a shortest path in a graph with $|V|$ vertices. This allows the algorithm to correctly calculate the shortest paths from the source vertex to all other vertices, **given that there are no negative-weight cycles.**

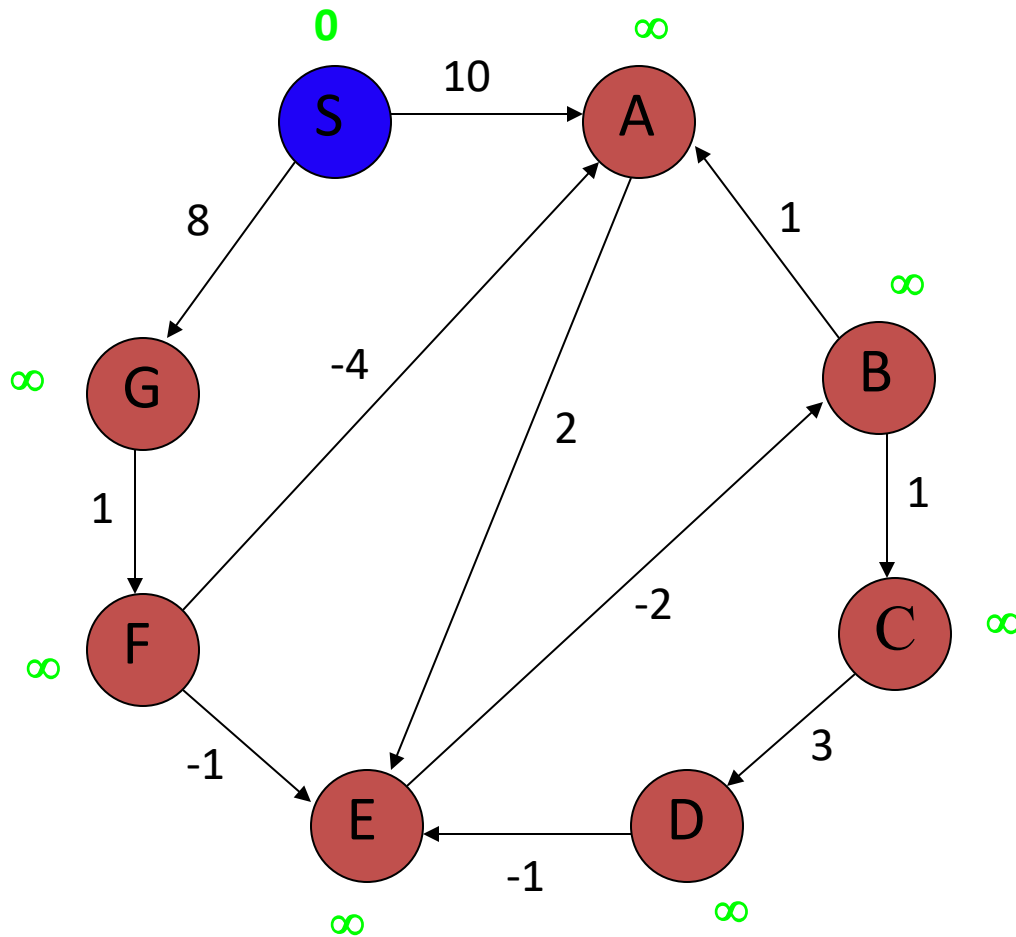
Bellman-Ford Algorithm with Negative Cycle Detection

- Initialize distance array `distTo[]` for each vertex `v` as **`distTo[v] = ∞`**, and **`distTo[s] = 0`** to source vertex `s`.
- Relax all **edges** $|V|-1$ times.
- Relax all the edges one more time i.e. the **N-th** time:
 - Case 1 (Negative cycle exists): if any edge can be further relaxed, i.e., for any **edge e**, if **`distTo[w] > distTo[v] + e.weight()`**
 - Case 2 (No Negative cycle) : case 1 fails for all the edges.

Time Complexity of Bellman-Ford Algorithm


- Time complexity for connected graph:
- Best Case: $O(|E|)$, when distance array after 1st and 2nd relaxation are same, we can simply stop further processing after one iteration
- Average Case: $O(|V|*|E|)$
- Worst Case: $O(|V|*|E|)$
 - If the graph is complete, the value of E becomes $O(|V|^2)$. So overall time complexity becomes $O(|V|^3)$

Bellman-Ford Algorithm Example 1



Iteration: 0

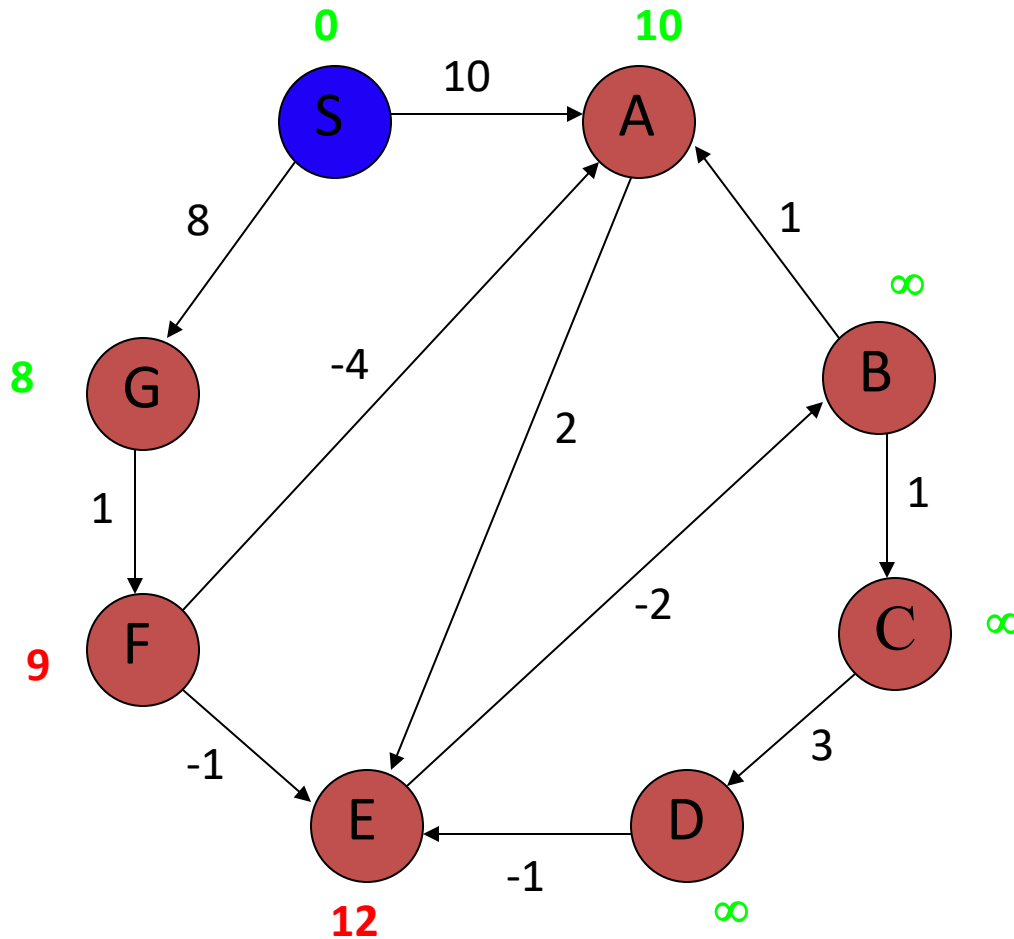
Bellman-Ford Algorithm Example 1



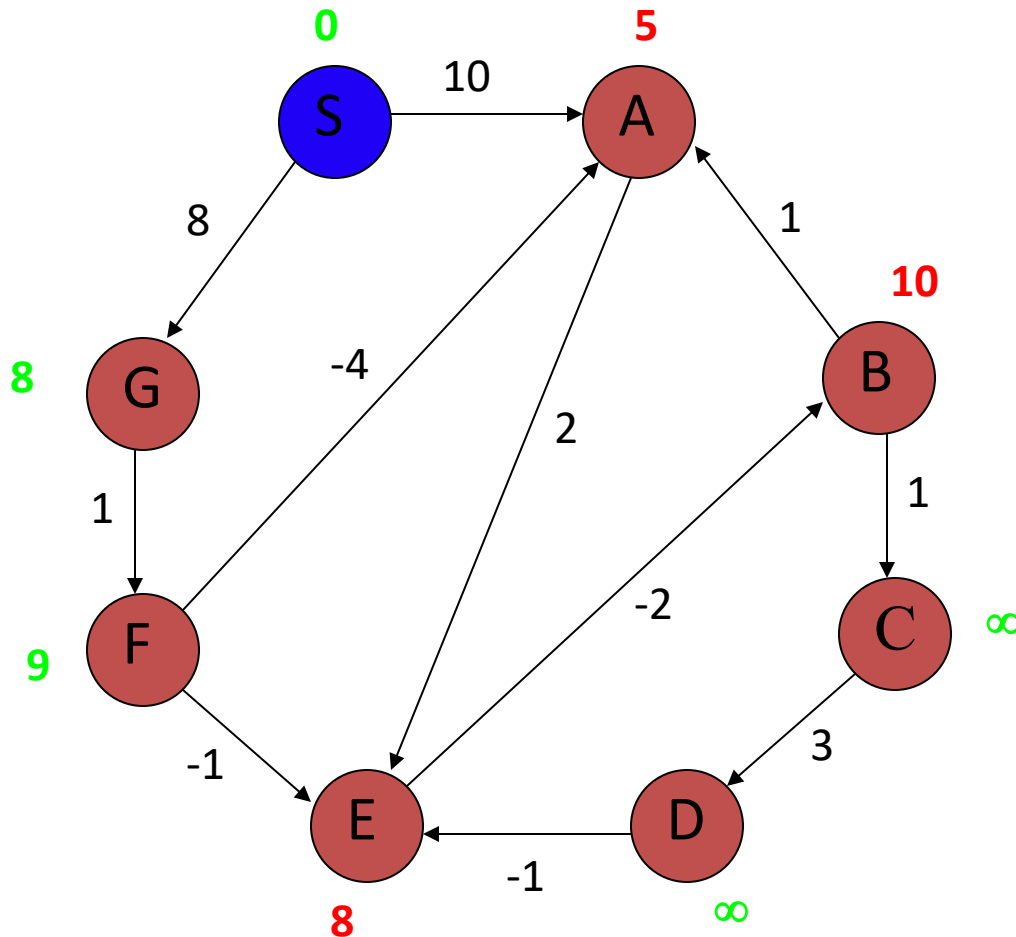
Iteration: 1

Bellman-Ford Algorithm Example 1

Iteration: 2



Bellman-Ford Algorithm Example 1

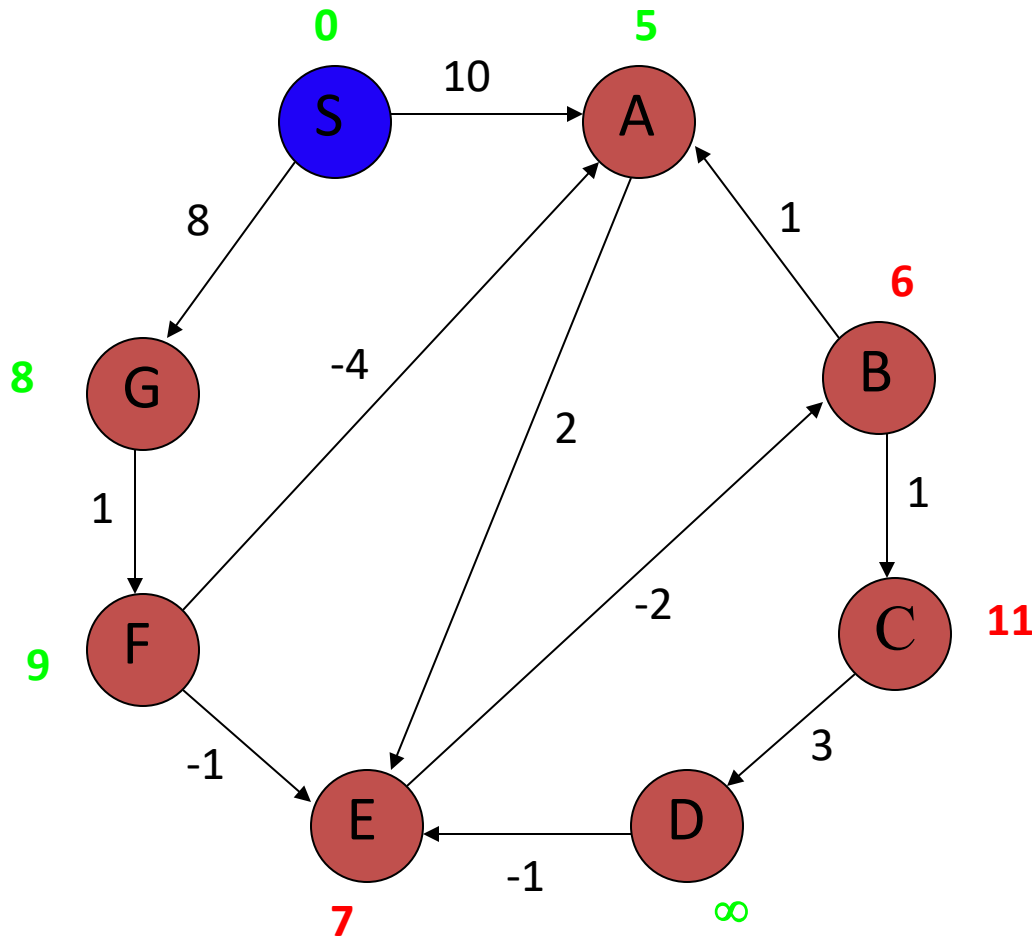


Iteration: 3

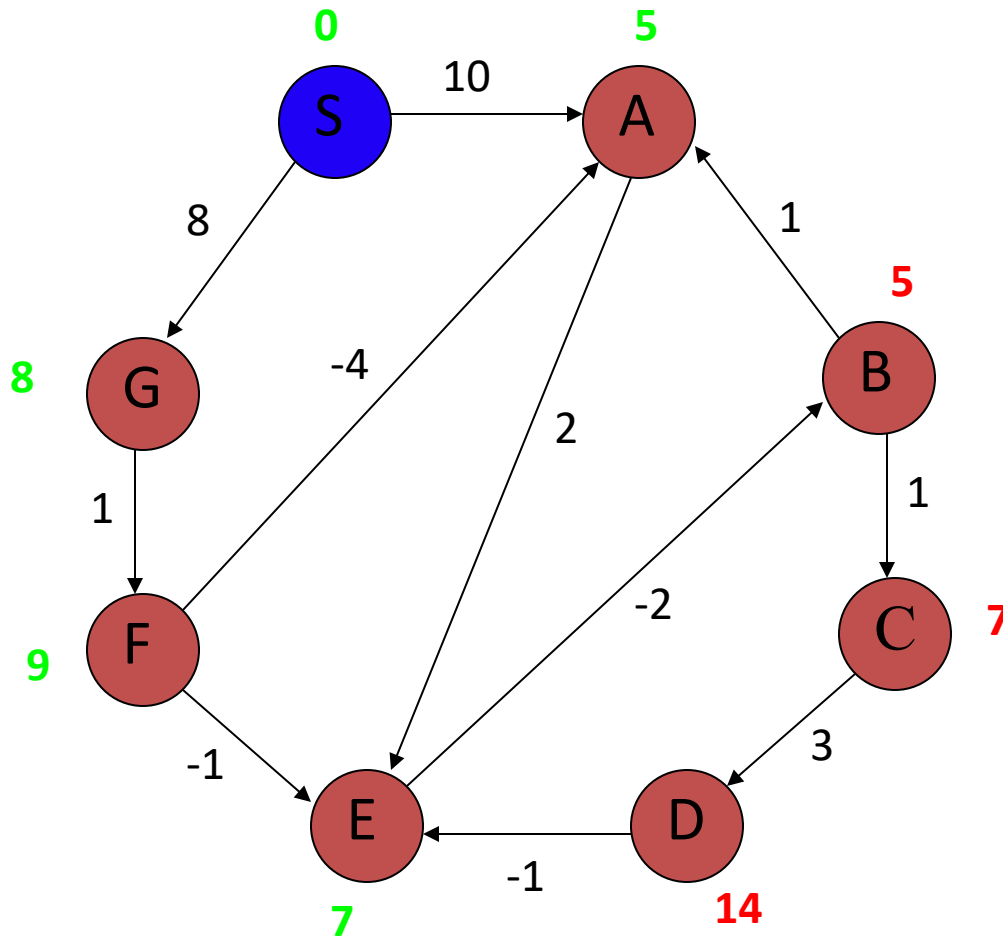
A has the correct
distance and path

Bellman-Ford Algorithm Example 1

Iteration: 4



Bellman-Ford Algorithm Example 1

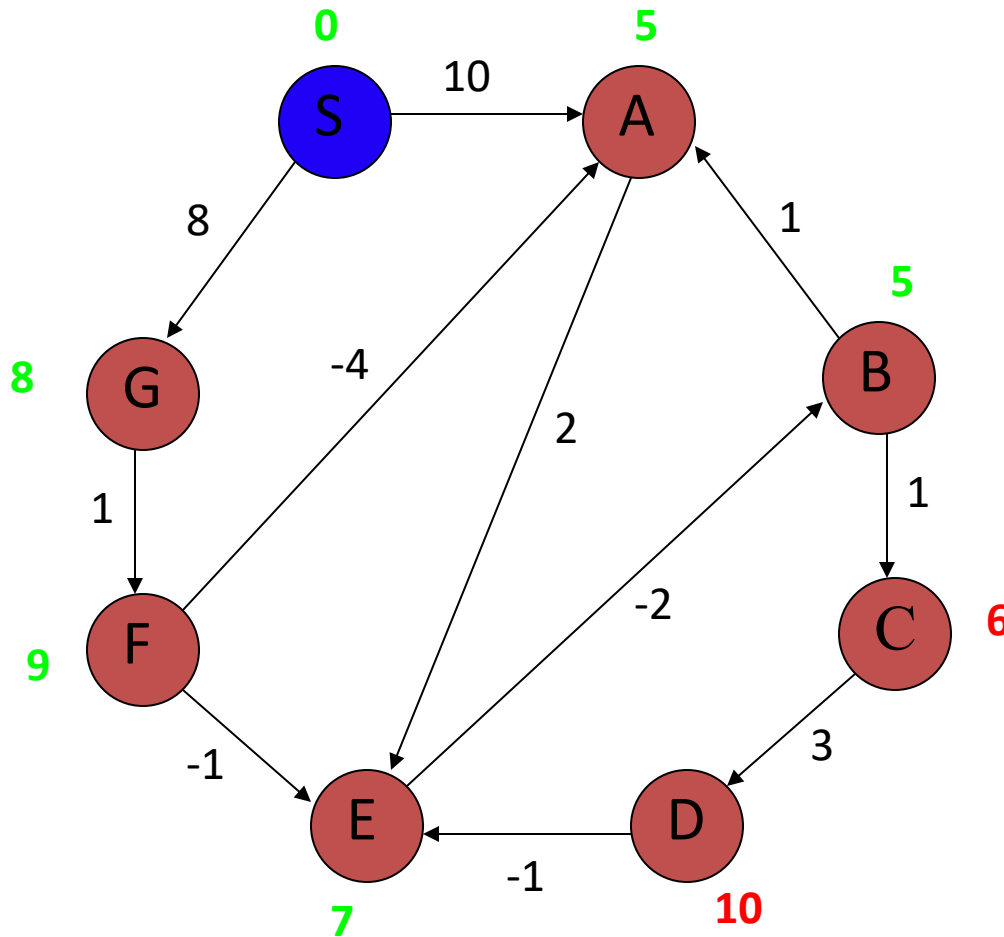


Iteration: 5

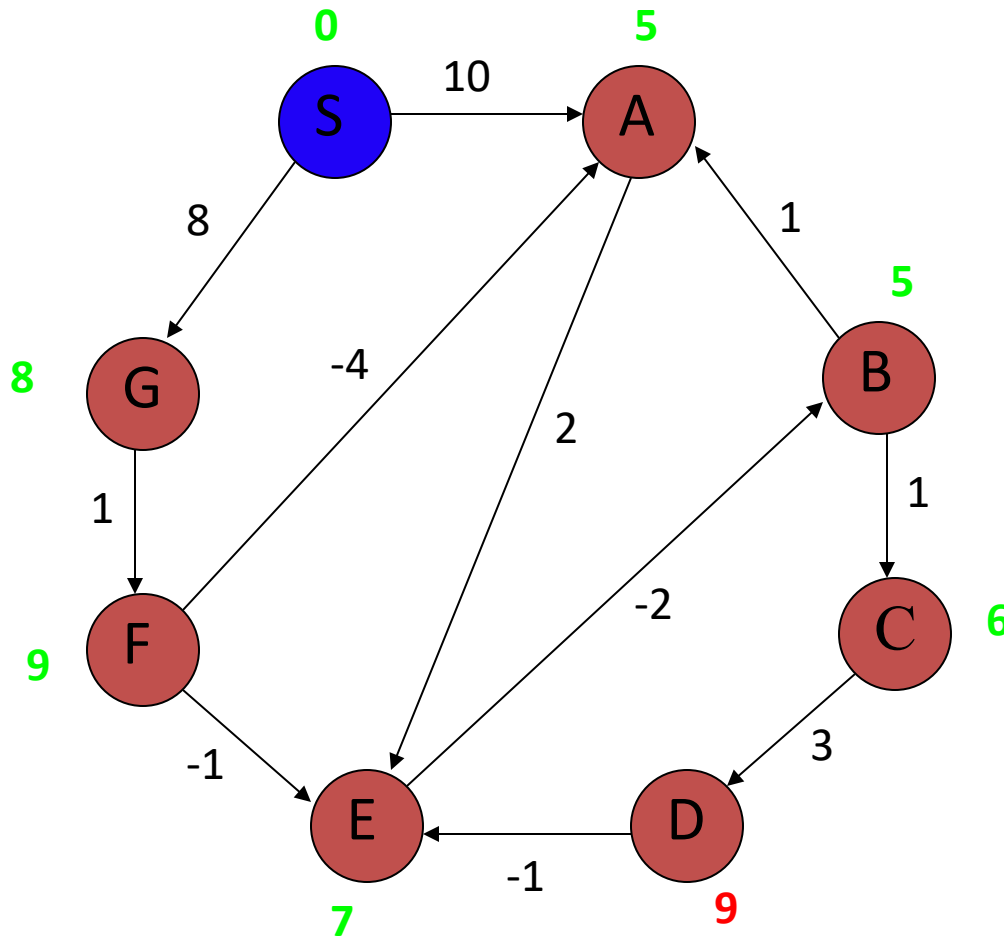
B has the correct
distance and path

Bellman-Ford Algorithm Example 1

Iteration: 6



Bellman-Ford Algorithm Example 1




Iteration: 7

D (and all other nodes) have the correct distance and path

Bellman-Ford Algorithm Example 2

Repeat $V - 1$ times: relax all E edges.



v	distTo[]		
0	∞	0	
1	∞	5	
2	∞	17	14
3	∞	20	17
4	∞	9	
5	∞	13	
6	∞	28	26 25
7	∞	8	

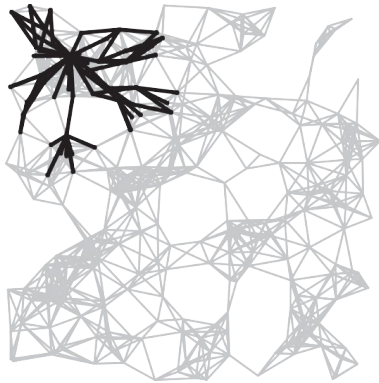
v	edgeTo[]		
0	-		
1	-	0	
2	-	1	5
3	-	1	2
4	-	0	
5	-	4	
6	-	2	5 2
7	-	0	

pass 1 pass 2 pass 3 (no further changes) pass 4-7 (no further changes)

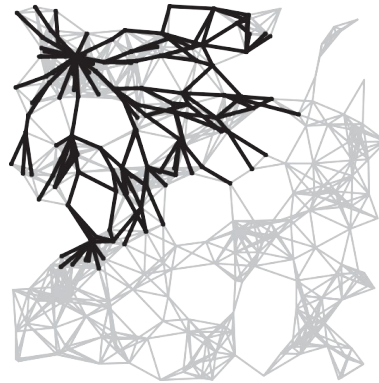
0→1 0→4 0→7 1→2 1→3 1→7 2→3 2→6 3→6 4→5 4→6 4→7 5→2 5→6 7→2 7→5

Bellman-Ford Algorithm Visualization

passes 4



7



10



13



SPT



Quiz

- Given a graph where all edges have positive weights, the shortest paths produced by Dijkstra and Bellman Ford algorithm may be different but path weight would always be same.
- ANS: True
- Dijkstra and Bellman-Ford both work fine for a graph with all positive weights, but they are different algorithms and may pick different edges for shortest paths.

Quiz

- Let G be a directed graph whose vertex set is the set of numbers from 1 to 100. There is an edge from a vertex i to a vertex j if either $j = i + 1$ or $j = 3i$. The minimum number of edges in a path in G from vertex 1 to vertex 100 is
- A. 4 B. 7 C. 23 D. 99
- ANS: 7
- The task is to find minimum number of edges in a path in G from vertex 1 to vertex 100 such that we can move to either $i+1$ or $3i$ from a vertex i .
- Since the task is to minimize number of edges, we would prefer to follow $3*i$. Let us follow multiple of 3. $1 \Rightarrow 3 \Rightarrow 9 \Rightarrow 27 \Rightarrow 81$, now we can't follow multiple of 3 anymore. So we will have to follow $i+1$. This solution gives a long path.
- What if we begin from end, and we reduce by 1 if the value is not multiple of 3, else we divide by 3. $100 \Rightarrow 99 \Rightarrow 33 \Rightarrow 11 \Rightarrow 10 \Rightarrow 9 \Rightarrow 3 \Rightarrow 1$
- So we need total 7 edges.